SUNDAY JUNE 12TH

	Olympia A	Olympia B	Vergina	Athens View	Pella	Horizon	Mycenae
08:00-09.00	**Continental Breakfast** (Foyer Olympia)						
09:00-10:30	**MobiDE Workshop**	**WebDB Workshop**	**NetDB Workshop**	NO SESSION	**DBSocial Workshop**	NO SESSION	**SWIM Workshop**
10:30-11:00	**Coffee Break** (Ballroom Foyer)						
11:00-12:30	**MobiDE Workshop**	**WebDB Workshop**	**NetDB Workshop**	NO SESSION	**DBSocial Workshop**	NO SESSION	**SWIM Workshop**
12:30-14:00	**Lunch Break** (on your own)						
14:00-15:30	**MobiDE Workshop**	**WebDB Workshop**	**NetDB Workshop**	NO SESSION	**DBSocial Workshop**	NO SESSION	**SWIM Workshop**
15:30-16:00	**Coffee Break** (Ballroom Foyer)						
16:00-17:00	**MobiDE Workshop**	**WebDB Workshop**	**NetDB Workshop**	NO SESSION	**DBSocial Workshop**	NO SESSION	**SWIM Workshop**
17:30-20:00	**PODS 30th Anniversary Colloquium** (Horizon)						
20:00-22:00	**PODS Reception** (Roof Garden)						

MONDAY JUNE 13TH

	Olympia	Vergina	Athens View	Pella	Horizon	Mycenae
08:00-08.30	**Continental Breakfast** (Foyer Olympia)					
08:45-09:00	**PODS Opening Ceremony** (Olympia)					
09:00-10:00	**PODS Keynote: : A Quest for Beauty and Wealth (or, Business Processes for Database Researchers)** *Daniel Deutch, Ben Gurion University & Tova Milo, Tel Aviv University* (Olympia)					
10:00-10:30	**Coffee Break** (Foyer Olympia)					
10:30-12:00	**PODS 1** Streaming and Sampling	NO SESSION	NO SESSION	**DaMoN Workshop**	NO SESSION	**DBTest Workshop**
12:00-13:30	**Lunch Break** (on your own)					
13:30-15:00	**PODS 2** Incomplete Information and Awards	NO SESSION	NO SESSION	**DaMoN Workshop**	NO SESSION	**DBTest Workshop**
15:00-15:30	**Coffee Break** (Foyer Olympia)					
15:30-17:00	**PODS 3** Index Structures and External Memory	NO SESSION	NO SESSION	**DaMoN Workshop**	NO SESSION	**DBTest Workshop**
17:00-17:30	**Coffee Break** (Foyer Olympia)					
17:30-19:00	NO SESSION	NO SESSION	NO SESSION	**DaMoN Workshop**	NO SESSION	**DBTest Workshop**
19:00-20:00	**Free time**					
20:00-22:30	**SIGMOD Reception and Undergraduate Poster Competition + IDAR ?????** (Roof Garden)					

TUESDAY JUNE 14TH

	Olympia A	Olympia B	Vergina	Athens View	Pella	Horizon	Mycenae
08:00-08.45	**Continental Breakfast** (Foyer Olympia)						
08:45-09:00	**A Greek-Style Welcome** (Olympia)						
09:00-10:00	**Keynote 1: Internet Scale Storage** *James Hamilton, VP & Distinguished Engineer, Amazon Web Services* (Olympia)						
10:00-10:30	**Coffee Break** (Foyer Olympia)						
10:30-11:30	**Research 1** Databases on New Hardware	**Research 2** Query Processing and Optimization	**Research 3** Schema Mapping and Data Integration	**Tutorial 1** Learning Statistical Models from Relational Data (Part 1)	**Industrial Keynote 1**	**PODS 4** Provenance	**Demo A** Ranking, the Web and Social Media
11:30-12:00					**Industrial 1** Data Management for Social Networking Applications		
12:00-13:30	**Lunch Break** (Ilissos/Amalia)						
13:30-15:00	**Research 4** Data on the Web	**Research 5** Data Privacy and Security	**Research 6** Data Consistency and Parallel DB	**Tutorial 1** Learning Statistical Models from Relational Data (Part 2)	**Industrial 2** Applying Hadoop	**PODS 5** Queries and Views	**Demo B & Programming Contest** Systems and Performance
15:00-15:30	**Coffee Break** (Foyer Olympia)						
15:30-16:30	**Research 7** Service Oriented Computing, Data Management in the Cloud	**Research 8** Spatial and Temporal Data Management	**Research 9** Shortest Paths and Sequence Data	**Tutorial 2** Datalog and Emerging Applications: An Interactive Tutorial	NO SESSION	**PODS Tutorial 1** Querying Semantic Web Data with SPARQL: State of the Art and Research Perspectives	**Demo C** Data Integration and Probabilistic Databases
16:30-17:30							
17:30-18:00	**Coffee Break** (Foyer Olympia)						
18:00-20:30	**New Researchers Symposium** (Olympia)						
20:30-23:00	**Microsoft Reception** (Roof Garden)						

WEDNESDAY JUNE 15TH

	Olympia A	Olympia B	Vergina	Athens View	Pella	Horizon	Mycenae
08:00-08.30	**Continental Breakfast** (Foyer Olympia)						
08:30-09:00	**Research 10** Data Provenance, Workflow and Cleaning	**Research 11** Information Extraction	**Research 12** Keyword Search and Ranked Queries	**Tutorial 3** Privacy-aware Data Management in Information Networks	**Industrial 3** Support for Business Analytics and Warehousing		**Demo D** User Support and Development Environments
09:00-10:00						**PODS Tutorial 2** Data Stream Analysis: Principles and Perspectives	
10:00-10:30	**Coffee Break** (Foyer Olympia)						
10:30-11:30	**Research 13** Stream and Complex Event Processing	**Research 14** Query Processing	**Research 15** Data Mining	**Tutorial 4** Web Data Management	**Industrial Keynote 2**	**PODS 6** Semi-structured Data and XML	**Demo B & Programming Contest** Systems and Performance
11:30-12:00					**Industrial 4** Business Analytics		
12:00-13:30	**Lunch Break** (Ilissos/Amalia)						
13:30-14:00	**Bussiness Meeting** (Olympia)						
14:00-15:45	Awards Presentations (Olympia)						
15:45-16:00	**Coffee Break** (Foyer Olympia)						
16:00-17:10	**Plenary I** Posters of Papers in Research Sessions 11 - 21 (Olympia)					**PODS 7** Rule-based Query Languages	**Demo A** Ranking, the Web and Social Media
17:10-18:00							
20:00-23:59	**Conference Banquet** (Island)						

THURSDAY JUNE 16TH

	Olympia A	Olympia B	Vergina	Athens View	Pella	Horizon	Mycenae
08:00-09.00	**Continental Breakfast** (Foyer Olympia)						
09:00-10.00	**Keynote 2: Managing Scientific Data** *Anastasia Ailamaki, EPFL* (Olympia)						
10:00-10:30	**Coffee Break** (Ballroom Foyer)						
10:30-12:00	**Research 16** Information Retrieval	**Research 17** Probabilistic and Uncertain Databases	**Research 18** Scalable Data Analytics	**Tutorial 5** Large-Scale Copy Detection		**Panel 1** Data Management Issues in Health and Medical Informatics	**Demo D** User Support and Development Environments
12:00-13:30	**Lunch Break** (on your own)						
13:30-15:00	**Research 19** Graph Management	**Research 20** OLAP and its Applications	**Research 21** Similarity Search and Queries	**Tutorial 6** Data Management Over Flash Memory	**Industrial 5** Dynamic Optimization and Unstructured Content	NO SESSION	**Demo C** Data Integration and Probabilistic Databases & Programming Contest Winners
15:00-15:30	**Coffee Break** (Ballroom Foyer)						
15:30-17:30	**Plenary II** Posters of Papers in Research Sessions 1 - 10 (Olympia)						

SESSION CONTENTS

SUNDAY 12TH 17:30 – 20:00

PODS 30th Anniversary Colloquium
Location: Horizon

SUNDAY 12TH 20:00 – 22:00

PODS Reception
Location: Roof Garden

MONDAY 13TH 08:45 – 10:00

PODS Opening and Keynote I
Location: Olympia

A Quest for Beauty and Wealth (or, Business Processes for Database Researchers)
Daniel Deutch (Ben Gurion University), Tova Milo (Tel Aviv University)

MONDAY 13TH 10:30 – 12:00

PODS Research 1: Streaming and Sampling
Location: Olympia

Get the Most Out of Your Sample: Optimal Unbiased Estimators Using Partial Information
Edith Cohen (AT&T Labs-Research), Haim Kaplan (Tel Aviv University)

Tight Bounds for Lp Samplers, Finding Duplicates in Streams, and Related Problems
Hossein Jowhari, Mert Sağlam (Simon Fraser University), Gábor Tardos (Rényi Institute of Mathematics & Simon Fraser University)

Pan-private Algorithms Via Statistics on Sketches
Darakhshan Mir, S. Muthukrishnan, Aleksandar Nikolov, Rebecca N. Wright (Rutgers University)

FIFO Indexes for Decomposable Problems
Cheng Sheng, Yufei Tao (Chinese University of Hong Kong)

MONDAY 13TH 13:30 – 15:00

PODS Research 2: Incomplete Information and Awards
Location: Olympia

ACM PODS Alberto O. Mendelzon Test-of-Time Award: Optimal Aggregation Algorithms for Middleware.
Ronald Fagin, Amnon Lotem, and Moni Naor.
Award committee: Peter Buneman (University of Edinburgh), Meral Ozsoyoglu (Case Western Reserve University), Jianwen Su (UC Santa Barbara)

Best Paper Award: Data Exchange Beyond Complete Data
Marcelo Arenas (PUC Chile), Jorge Pérez (Universidad de Chile), Juan Reutter (University of Edinburgh)

Regular Paper: Incomplete Information and Certain Answers in General Data Models
Leonid Libkin (University of Edinburgh)

Regular Paper: Determining the Currency of Data
Wenfei Fan (University of Edinburgh & Harbin Institute of Technology), Floris Geerts (University of Edinburgh), Jef Wijsen (Université de Mons)

MONDAY 13TH 15:30 – 17:00

PODS Research 3: Index Structures and External Memory
Location: Olympia

New Results on Two-dimensional Orthogonal Range Aggregation in External Memory
Cheng Sheng, Yufei Tao (Chinese University of Hong Kong)

On Finding Skylines in External Memory
Cheng Sheng , Yufei Tao (Chinese University of Hong Kong)

Beyond Simple Aggregates: Indexing for Summary Queries
Zhewei Wei, Ke Yi (Hong Kong University of Science and Technology)

Space-Efficient Substring Occurrence Estimation
Alessio Orlandi (University of Pisa), Rossano Venturini (ISTI-CNR)

MONDAY 13TH 20:00 – 22:30

SIGMOD Reception and Undergraduate Poster Competition
Location: Roof Garden

TUESDAY 14TH 08:45 – 10:00

SIGMOD Welcome and Keynote I
Location: Olympia

Internet-Scale Storage
James Hamilton (Amazon Web Services)

TUESDAY 14TH 10:30 – 12:00

SIGMOD Research 1: Databases on New Hardware
Location: Olympia A

LazyFTL: A Page-level Flash Translation Layer Optimized for NAND Flash Memory
Dongzhe Ma (Tsinghua University), Jianhua Feng (Tsinghua University), Guoliang Li (Tsinghua University)

Operation-Aware Buffer Management in Flash-based Systems
Yanfei Lv (Peking University); Bin Cui (Pku); Bingsheng He (Nanyang Technological Universi); Xuexuan Chen (Peking University)

SkimpyStash: RAM Space Skimpy Key-Value Store on Flash-based Storage
Biplob Debnath (EMC Corporation); Sudipta Sengupta (Microsoft Research); Jin Li (Microsoft Research)

Design and Evaluation of Main Memory Hash Join Algorithms for Multi-core CPUs
Spyros Blanas (Univ. of Wisconsin); Yinan Li (Univ. of Wisconsin); Jignesh Patel (Univ. of Wisconsin)

SIGMOD Research 2: Query Processing and Optimization
Location: Olympia B

Query Optimization Techniques for Partitioned Tables
Herodotos Herodotou (Duke University); Nedyalko Borisov (Duke University); Shivnath Babu (Duke University)

CrowdDB: Answering Queries with Crowdsourcing
Michael Franklin (UC Berkeley); Donald Kossmann (ETH Zurich); Tim Kraska (UC Berkeley); Sukriti Ramesh (ETH Zurich); Reynold Xin (UC Berkeley)

Skyline Query Processing over Joins
Akrivi Vlachou (NTNU); Christos Doulkeridis (NTNU); Neoklis Polyzotis (UC - Santa Cruz)

Efficient Parallel Skyline Processing using Hyperplane Projections
Henning Koehler (University of Queensland); Jing Yang; Xiaofang Zhou (University of Queensland)

SIGMOD Research 3: Schema Mapping and Data Integration
Location: Vergina

Scalable Query Rewriting: A Graph-Based Approach
George Konstantinidis (ISI/USC); Jose Luis Ambite (ISI/USC)

Automatic Discovery of Attributes in Relational Databases
Meihui Zhang (National University of Singapore); Marios Hadjieleftheriou (AT&T Labs - Research); Beng Chin Ooi (National University of Singapore); Cecilia Procopiuc (AT&T Labs - Research); Divesh Srivastava (AT&T Labs - Research)

Leveraging Query Logs for Schema Mapping Generation in U-MAP
Hazem Elmeleegy (Purdue University); Ahmed Elmagarmid (Qatar Computing Research Institute); Jaewoo Lee (Purdue University)

Designing and Refining Schema Mappings via Data Examples
Bogdan Alexe (UC Santa Cruz); Balder ten Cate (UC Santa Cruz); Phokion Kolaitis (UC Santa Cruz); Wang-Chiew Tan (IBM Research - Almaden and UC Santa Cruz)

SIGMOD Tutorial 1: Learning Statistical Models from Relational Data (Part 1)
Location: Athens View
Presenters: Lise Getoor (Univ. of Maryland); Lilyana Mihalkova (Univ. of Maryland)

SIGMOD Industrial Keynote and Session 1: Data Management for Social Networking Applications
Location: Pella

SIGMOD Industrial Keynote 1: Not Available
Michael Abbot (Twitter)

Bistro Data Feed Management System
Vladislav Shkapenyuk (AT&T Labs - Research), Theodore Johnson (AT&T Labs - Research), Divesh Srivastava (AT&T Labs - Research)

PODS Research 4: Provenance
Location: Horizon

Provenance for Aggregate Queries
Yael Amsterdamer (Tel Aviv University and University of Pennsylvania), Daniel Deutch (Ben Gurion University and University of Pennsylvania), Val Tannen (University of Pennsylvania)

On Provenance Minimization
Yael Amsterdamer (Tel Aviv University and University of Pennsylvania), Daniel Deutch (Ben Gurion University and University of Pennsylvania), Tova Milo (Tel Aviv University), Val Tannen (University of Pennsylvania)

On the Complexity of Privacy-Preserving Complex Event Processing
Yeye He, Siddharth Barman (University of Wisconsin-Madison), Di Wang (Worcester Polytechnic Institute), Jeffrey F. Naughton (University of Wisconsin-Madison)

Provenance Views for Module Privacy
Susan B. Davidson, Sanjeev Khanna (University of Pennsylvania), Tova Milo (Tel Aviv University), Debmalya Panigrahi (Massachusetts Institute of Technology), Sudeepa Roy (University of Pennsylvania)

SIGMOD Demo A: Ranking, the Web and Social Media
Location: Mycenae

SkylineSearch: Semantic Ranking and Result Visualization for PubMed
Julia Stoyanovich (University of Pennsylvania), Mayur Lodha (Columbia University), William Mee (Columbia University), Kenneth A. Ross (Columbia University)

A Cross-Service Travel Engine for Trip Planning
Gang Chen (Zhejiang University), Chen Liu (National University of Singapore), Meiyu Lu (National University of Singapore), Beng Chin Ooi (National University of Singapore), Shanshan Ying (National University of Singapore), Anthony Tung (National University of Singapore), Dongxiang Zhang (National University of Singapore), Meihui Zhang (National University of Singapore)

WINACS: Construction and Analysis of Web-Based Computer Science Information Networks
Tim Weninger (University of Illinois Urbana-Champaign), Marina Danilevsky (University of Illinois Urbana-Champaign), Fabio Fumarola (Universita), Joshua Hailpern (University of Illinois Urbana-Champaign), Jiawei Han (University of Illinois Urbana-Champaign), Thomas J. Johnston (University of Illinois Urbana-Champaign), Surya Kallumadi (Kansas State University), Hyungsul Kim (University of Illinois Urbana-Champaign), Zhijin Li (University of Illinois Urbana-Champaign), David McCloskey (University of Illinois Urbana-Champaign), Yizhou Sun (University of Illinois Urbana-Champaign), Nathan E. TeGrotenhuis

(Whitworth University), Chi Wang (University of Illinois Urbana-Champaign), Xiao Yu (University of Illinois Urbana-Champaign)

Tweets as Data: Demonstration of TweeQL and TwitInfo
Adam Marcus (MIT CSAIL), Michael S. Bernstein (MIT CSAIL), Osama Badar (MIT CSAIL), David R. Karger (MIT CSAIL), Samuel Madden (MIT CSAIL), Robert C. Miller (rcm@csail.mit.edu)

MOBIES: Mobile-Interface Enhancement Service for Hidden Web Database
Xin Jin (George Washington University), Aditya Mone (University of Texas at Arlington), Nan Zhang (George Washington University), Gautam Das (University of Texas at Arlington)

Search Computing: Multi-domain Search on Ranked Data
Alessandro Bozzon (Politecnico di Milano), Daniele Braga (Politecnico di Milano), Marco Brambilla (Politecnico di Milano), Stefano Ceri (Politecnico di Milano), Francesco Corcoglioniti (Politecnico di Milano), Piero Fraternali (Politecnico di Milano), Salvatore Vadacca (Politecnico di Milano)

EnBlogue -- Emergent Topic Detection in Web 2.0 Streams
Foteini Alvanaki (Saarland University), Michel Sebastian (Saarland University), Krithi Ramamritham (IIT Bombay), Gerhard Weikum (Max-Planck Institute Informatics)

NOAM: News Outlets Analysis and Monitoring System
Ilias Flaounas (University of Bristol), Omar Ali (University of Bristol), Marco Turchi (European Commission), Tristan Snowsill (University of Bristol), Florent Nicart (Université de Rouen), Tijl De Bie (University of Bristol), Nello Cristianini (University of Bristol)

TUESDAY 14TH 13:30 – 15:00

SIGMOD Research 4: Data on the Web
Location: Olympia A

Apples and Oranges: A Comparison of RDF Benchmarks and Real RDF Datasets
Songyun Duan (IBM Research - Thomas J. Watson Research Ctr.), Anastasios Kementsietsidis (IBM T.J. Watson), Kavitha Srinivas (IBM T. J. Watson), Octavian Udrea (IBM T. J. Watson)

Efficient Query Answering in Probabilistic RDF Graphs
Xiang Lian (HKUST), Lei Chen (Hong Kong University of Science and Technology)

Facet Discovery for Structured Web Search: A Query-log Mining Approach
Jeffrey Pound (University of Waterloo), Stelios Paparizos (Microsoft Research), Panayiotis Tsaparas

Schema-As-You-Go: On Probabilistic Tagging and Querying of Wide Tables
Meiyu Lu (NUS), Bing Tian Dai (NUS), Anthony Tung, Divyakant Agrawal (UC Santa Barbara)

SIGMOD Research 5: Data Privacy and Security
Location: Olympia B

No Free Lunch in Data Privacy
Daniel Kifer (Penn State), Ashwin Machanavajjhala (Yahoo)

TrustedDB: A Trusted Hardware based Database with Privacy and Data Confidentiality
Sumeet Bajaj (Stony Brook), Radu Sion (SUNY Stony Brook)

Differentially Private Data Cubes: Optimizing Noise Sources and Consistency
Bolin Ding (UIUC), Marianne Winslett (UIUC), Jiawei Han (UIUC); Zhenhui Li (UIUC)

iReduct: Differential Privacy with Reduced Relative Errors
Xiaokui Xiao (Nanyang Technological Univ), Gabriel Bender (Cornell University), Michael Hay (Cornell University), Johannes Gehrke (Cornell University)

SIGMOD Research 6: Data Consistency and Parallel DB
Location: Vergina

A Latency and Fault-Tolerance Optimizer for Online Parallel Query Plans
Prasang Upadhyaya (University of Washington), YongChul Kwon (University of Washington), Magdalena Balazinska (University of Washington)

ArrayStore: A Storage Manager for Complex Parallel Array Processing
Emad Soroush (Univ. of Washington), Magdalena Balazinska (Univ. of Washington), Daniel Wang (SLAC)

Fast Checkpoint Recovery Algorithms for Frequently Consistent Applications

Tuan Cao (Cornell University), Marcos Vaz Salles (Cornell University), Benjamin Sowell (Cornell University), Yao Yue (Cornell University), Johannes Gehrke (Cornell University), Alan Demers (Cornell University), Walker White (Cornell University)

Warding off the Dangers of Data Corruption with Amulet
Nedyalko Borisov (Duke University), Shivnath Babu (Duke University), Nagapramod Mandagere (IBM Almaden Research Center), Sandeep Uttamchandani (IBM Almaden Research Center)

SIGMOD Tutorial 1: Learning Statistical Models from Relational Data (Part 2)
Location: Athens View
Presenters: Lise Getoor (Univ. of Maryland); Lilyana Mihalkova (Univ. of Maryland)

SIGMOD Industrial 2: Applying Hadoop
Location: Pella

Apache Hadoop Goes Realtime at Facebook
Dhruba Borthakur (facebook), Jonathan Gray (facebook), Joydeep Sen Sarma (facebook), Kannan Muthukkaruppan (facebook), Nicolas Spiegelberg (facebook), Hairong Kuang (facebook), Karthik Ranganathan (facebook), Dmytro Molkov (facebook), Aravind Menon (facebook), Samuel Rash (facebook), Rodrigo Schmidt (facebook), Amitanand Aiyer (facebook)

Nova: Continuous Pig/Hadoop Workflows
Christopher Olston (Yahoo! Research), Greg Chiou (Yahoo!), Laukik Chitnis (Yahoo!), Francis Liu (Yahoo!), Yiping Han (Yahoo!), Mattias Larsson (Yahoo!), Andreas Neumann (Yahoo!), Vellanki B. N. Rao (Yahoo!), Vijayanand Sankarasubramanian (Yahoo!), Siddharth Seth (Yahoo!), Chao Tian (Yahoo!), Topher ZiCornell (Yahoo!), Xiaodan Wang (Johns Hopkins University)

A Hadoop Based Distributed Loading Approach to Parallel Data Warehouses
Yu Xu (Teradata), Pekka Kostamaa (Teradata), Yan Qi (Teradata), Jian Wen (UC Riverside), Kevin Keliang Zhao (UC San Diego)

A Batch of PNUTS: Experiences Connecting Cloud Batch and Serving Systems
Adam E Silberstein (Yahoo! Research), Russell Sears (Yahoo! Research), Wenchao Zhou (Univeristy of Pennsylvania), Brian Frank Cooper (Yahoo! Research)

PODS Research 5: Queries and Views
Location: Horizon

Querying Graph Patterns
Pablo Barcelo (Universidad de Chile), Leonid Libkin, Juan L. Reutter (University of Edinburgh)

Maximizing Conjunctive Views in Deletion Propagation
Benny Kimelfeld, Jan Vondrák, Ryan Williams (IBM Research - Almaden)

Determining Relevance of Accesses at Runtime
Michael Benedikt, Georg Gottlob (University of Oxford), Pierre Senellart (Télécom ParisTech)

Parallel Evaluation of Conjunctive Queries
Paraschos Koutris, Dan Suciu (University of Washington)

SIGMOD Demo B & Programming Contest: Systems and Performance
Location: Mycenae

One-pass Data Mining Algorithms in a DBMS with UDFs
Carlos Ordonez (University of Houston), Sasi K Pitchaimalai (University of Houston)

Inspector Gadget: A Framework for Custom Monitoring and Debugging of Distributed Dataflows
Christopher Olston (Yahoo! Research), Benjamin Reed (Yahoo! Research)

RAFT at Work: Speeding-Up MapReduce Applications under Task and Node Failures
Jorge-Arnulfo Quiané-Ruiz (Saarland University), Christoph Pinkel (Saarland University), Jörg Schad (Saarland University), Jens Dittrich (Saarland University)

WattDB: An Energy-Proportional Cluster of Wimpy Nodes
Daniel Schall (TU Kaiserslautern), Volker Hudlet (TU Kaiserslautern)

BRRL: A Recovery Library for Main-Memory Applications in the Cloud
Tuan Cao (Cornell University), Benjamin Sowell (Cornell University), Marcos Vaz Salles (Cornell University), Alan Demers (Cornell University), Johannes Gehrke (Cornell University)

A Data-oriented Transaction Execution Engine and Supporting Tools
Ippokratis Pandis (Carnegie Mellon University), Pinar Tozun (Ecole Polytechnique Federale de Lausanne), Miguel Branco (Ecole Polytechnique Federale de Lausanne), Dimitris Karampinas (University of Patras),

Danica Porobic (Ecole Polytechnique Federale de Lausanne), Ryan Johnson (University of Toronto), Anastasia Ailamaki (Ecole Polytechnique Federale de Lausanne)

iGraph in Action: Performance Analysis of Disk-Based Graph Indexing Techniques
Wook-Shin Han (Kyungpook National University), Minh-Duc Pham (Kyungpook National University), Jinsoo Lee (Kyungpook National University), Romans Kasperovics (Kyungpook National University), Jeffrey Xu Yu (Chinese University of Hong Kong)

StreamRec: A Real-Time Recommender System
Badrish Chandramouli (Microsoft Research), Justin J Levandoski (University of Minnesota), Ahmed Eldawy (University of Minnesota), Mohamed F Mokbel (University of Minnesota)

TUESDAY 14TH 15:30 – 17:30

SIGMOD Research 7: Service Oriented Computing, Data Management in the Cloud
Location: Olympia A

Schedule Optimization for Data Processing Flows on the Cloud
Herald Kllapi (University Of Athens), Eva Sitaridi (University of Athens), Manolis Tsangaris (University of Athens), Yannis Ioannidis (University of Athens)

Zephyr: Live Migration in Shared Nothing Databases for Elastic Cloud Platforms
Aaron Elmore (UC Santa Barbara), Sudipto Das (UC Santa Barbara), Divyakant Agrawal (UC Santa Barbara), Amr El Abbadi (UC Santa Barbara)

Workload-Aware Database Monitoring and Consolidation
Carlo Curino (MIT), Evan Jones (MIT), Sam Madden (MIT), Hari Balakrishnan (MIT)

Predicting Cost Amortization for Query Services
Verena Kantere (Cyprus University of Technology), Debabrata Dash (ArcSight), Georgios Gratsias (ELCA Informatique SA), Anastasia Ailamaki (EPFL)

Performance Prediction for Concurrent Database Workloads
Jennie Duggan (Brown University), Ugur Cetintemel (Brown University), Olga Papaemmanouil (Brandeis University), Eli Upfal (Brown University)

SIGMOD Research 8: Spatial and Temporal Data Management
Location: Olympia B

Reverse Spatial and Textual k Nearest Neighbor Search
Jiaheng Lu (Renmin University of China), Ying Lu (Renmin University of China), Gao Cong (Nanyang Technological University)

Location-Aware Type Ahead Search on Spatial Databases : Semantics and Efficiency
Senjuti Basu Roy (UT Arlington), Kaushik Chakrabarti (Microsoft Research)

Collective Spatial Keyword Querying
Xin Cao (Singapore NTU), Gao Cong (Nanyang Technological University), Christian Jensen (Aarhus University), Beng Chin Ooi (National University of Singapore)

Finding Semantics in Time Series
Peng Wang (Fudan University), Haixun Wang (Microsoft), Wei Wang (Fudan University)

Querying Contract Databases Based on Temporal Behavior
Elio Damaggio (UCSD); Alin Deutsch (UCSD); Dayou Zhou (UCSD)

SIGMOD Research 9: Shortest Paths and Sequence Data
Location: Vergina

Neighborhood-Privacy Protected Shortest Distance Computing in Cloud
Jun Gao (Peking University), Jeffrey Xu Yu (The Chinese University of Hong Kong), Ruoming Jin (Kent State University), Jiashuai Zhou, Tengjiao Wang, Dongqing Yang

On k-skip Shortest Paths
Yufei Tao (Chinese Univ. of Hong Kong), Cheng Sheng (The Chinese University of Hong Kong), JIan Pei (Simon Fraser Univ.)

Finding Shortest Path on Land Surface
Lian Liu (HKUST); Raymond Chi-Wing Wong (The Hong Kong University of Sc)

WHAM: A High-throughput Sequence Alignment Method

Yinan Li (Univ. of Wisconsin), Allison Terrell (Univ. of Wisconsin), Jignesh Patel (Univ. of Wisconsin)

A New Approach for Processing Ranked Subsequence Matching Based on Ranked Union
Wook-Shin Han (Kyungpook National University), Jinsoo Lee (Kyungpook National University), Yang-Sae Moon, Seung-won Hwang, Hwanjo Yu (POSTECH)

SIGMOD Tutorial 2: Datalog and Emerging Applications: An Interactive Tutorial
Location: Athens View
Presenters: Shan Shan Huang (MIT), Todd J. Green (UC Davis), Boon Thau Loo (Univ. of Pennsylvania)

PODS Tutorial 1: Querying Semantic Web Data with SPARQL: State of the Art and Research
Location: Horizon
Presenters: Marcelo Arenas (PUC Chile), Jorge Pérez (Universidad de Chile)

SIGMOD Demo C: Data Integration and Probabilistic Databases
Location: Mycenae

Pay-As-You-Go Mapping Selection in Dataspaces
Cornelia Hedeler (The University of Manchester), Khalid Belhajjame (The University of Manchester), Norman W Paton (The University of Manchester), Alvaro A A Fernandes (The University of Manchester), Suzanne M Embury (The University of Manchester), Lu Mao (The University of Manchester), Chenjuan Guo (The University of Manchester)

Exelixis: Evolving Ontology-Based Data Integration System
Haridimos Kondylakis (FORTH-ICS), Dimitris Plexousakis (FORTH-ICS)

U-MAP: A System for Usage-Based Schema Matching and Mapping
Hazem Elmeleegy (AT&T Labs - Research), Jaewoo Lee (Purdue University), El Kindi Rezig (Purdue University), Mourad Ouzzani (Purdue University), Ahmed Elmagarmid (Qatar Computing Research Institute, Qatar Foundation)

The SystemT IDE: An Integrated Development Environment for Information Extraction Rules
Laura Chiticariu (IBM Research - Almaden), Vivian Chu (IBM research - Almaden), Sajib Dasgupta (IBM Research - Almaden), Thilo W Goetz (IBM Software - Germany), Howard Ho (IBM Research - Almaden), Rajasekar Krishnamurthy (IBM Research - Almaden), Alexander Lang (IBM Software - Germany), Yunyao Li (IBM Research - Almaden), Bin Liu (University of Michigan), Sriram Raghavan (IBM Research - India), Frederick R Reiss (IBM Research - Almaden), Shivakumar Vaithyanathan (IBM Research - Almaden), Huaiyu Zhu (IBM Research - Almaden)

ProApproX: A Lightweight Approximation Query Processor over Probabilistic Trees
Pierre Senellart (Institut Télécom; Télécom ParisTech), Asma Souihli (Institut Télécom; Télécom ParisTech)

SPROUT²: A Squared Query Engine for Uncertain Web Data
Robert Fink (University of Oxford), Andrew Hogue (Google Inc.), Dan Olteanu (University of Oxford), Swaroop Rath (University of Oxford)

Fuzzy Prophet: Parameter exploration in uncertain enterprise scenarios
Oliver A Kennedy (EPFL), Steve Lee (Microsoft Corporation), Charles Loboz (Microsoft Corporation), Slawek Smyl (Microsoft Corporation), Suman Nath (Microsoft Research)

LinkDB: A Probabilistic Linkage Database System
Ekaterini Ioannou (Technical University of Crete), Wolfgang Nejdl (L3S Research Center), Claudia Niederée (L3S Research Center), Yannis Velegrakis (University of Trento)

TUESDAY 14TH 18:00 – 20:30

SIGMOD New Researchers Symposium
Location: Olympia

TUESDAY 14TH 20:30 – 23:00

Microsoft Reception
Location: Roof Garden

WEDNESDAY 15TH 08:30 – 10:00

SIGMOD Research 10: Data Provenance, Workflow and Cleaning
Location: Olympia A

Interaction between Record Matching and Data Repairing
Wenfei Fan (University of Edinburgh), Jianzhong Li (Harbin Institute of Technology), Shuai Ma (Beihang University), Nan Tang (University of Edinburgh), Wenyuan Yu (University of Edinburgh)

We Challenge You to Certify Your Updates
Su Chen (National University of Singapore), Xin Luna Dong (AT&T Labs-Research), Laks V.S. Lakshmanan (University of British Columbia), Divesh Srivastava (AT&T Labs-Research)

Labeling Recursive Workflow Executions On-the-Fly
Zhuowei Bao (University of Pennsylvania), Susan B. Davidson (University of Pennsylvania), Tova Milo (Tel Aviv University)

Tracing Data Errors with View-Conditioned Causality
Alexandra Meliou (University of Washington), Wolfgang Gatterbauer (University of Washington), Suman Nath (Microsoft Research), Dan Suciu (University of Washington)

SIGMOD Research 11: Information Extraction
Location: Olympia B

Hybrid In-Database Inference for Declarative Information Extraction
Daisy Zhe Wang (University of California, Berkeley), Michael J. Franklin (University of California, Berkeley), Minos Garofalakis (Technical University of Crete), Joseph M. Hellerstein (University of California, Berkeley), Michael L. Wick (University of Massachusetts, Amherst)

Faerie: Efficient Filtering Algorithms for Approximate Dictionary-based Entity Extraction
Guoliang Li (Tsinghua University), Dong Deng (Tsinghua University), Jianhua Feng (Tsinghua University)

Joint Unsupervised Structure Discovery and Information Extraction
Eli Cortez (Universidade Federal do Amazonas), Daniel Oliveira (Universidade Federal do Amazonas), Altigran S. da Silva (Universidade Federal do Amazonas), Edleno S. de Moura (Universidade Federal do Amazonas), Alberto H. F. Laender (Universidade Federal de Minas Gerais)

Attribute Domain Discovery for Hidden Web Databases
Xin Jin (George Washington University), Nan Zhang (George Washington University), Gautam Das (University of Texas at Arlington)

SIGMOD Research 12: Keyword Search and Ranked Queries
Location: Vergina

Keyword Search over Relational Databases: A Metadata Approach
Sonia Bergamaschi (University of Modena and Reggio Emilia, Italy), Elton Domnori (University of Modena and Reggio Emilia, Italy), Francesco Guerra (University of Modena and Reggio Emilia, Italy), Raquel Trillo Lado (University of Zaragoza,), Yannis Velegrakis (University of Trento)

Sharing Work in Keyword Search over Databases
Marie Jacob (University Of Pennsylvania), Zachary Ives (University Of Pennsylvania)

Nearest Keyword Search in XML Documents
Yufei Tao (Chinese University of Hong Kong), Stavros Papadopoulos (Chinese University of Hong Kong), Cheng Sheng (Chinese University of Hong Kong), Kostas Stefanidis (Chinese University of Hong Kong)

Efficient and Generic Evaluation of Ranked Queries
Wen Jin (Independent Consultant), Jignesh M Patel (University of Wisconsin--Madison)

SIGMOD Industrial 3: Support for Business Analytics and Warehousing
Location: Pella

Emerging Trends in the Enterprise Data Analytics
Fatma Özcan (IBM), David Hoa (IBM), Kevin S. Beyer (IBM), Andrey Balmin (IBM), Chuan Jie Liu (IBM), Yu Li (IBM)

Efficient Processing of Data Warehousing Queries in a Split Execution Environment
Kamil Bajda-Pawlikowski (Hadapt Inc. & Yale University), Daniel J Abadi (Hadapt Inc. & Yale University), Avi Silberschatz (Yale University), Erik Paulson (University of Wisconsin-Madison)

SQL Server Column Store Indexes

Per-Åke Larson (Microsoft), Cipri Clinciu (Microsoft), Eric N Hanson (Microsoft), Artem Oks (Microsoft), Susan L Price (Microsoft), Srikumar Rangarajan (Microsoft), Aleksandras Surna (Microsoft), Qingqing Zhou (Microsoft)

An Analytic Data Engine for Visualization in Tableau
Richard Wesley (Tableau Software), Matthew Eldridge (Tableau Software), Pawel T Terlecki (Tableau Software)

SIGMOD Tutorial 3: Privacy-aware Data Management in Information Networks
Location: Olympia A
Presenters: Michael Hay (Cornell University), Kun Liu (Yahoo! Labs), Gerome Miklau (University of Massachusetts Amherst), Jian Pei (Simon Fraser University), Evimaria Terzi (Boston University)

PODS Tutorial 2: Streams and Complex Event Processing
Location: Horizon
Presenter: S. Muthukrishnan (Rutgers University)

SIGMOD Demo D: User Support and Development Environments
Location: Mycenae

CONFLuEnCE: CONtinuous workFLow ExeCution Engine
Panayiotis Neophytou (University of Pittsburgh), Panos K. Chrysanthis (University of Pittsburgh), Alexandros Labrinidis (University of Pittsburgh)

Demonstration of Qurk: A Query Processor for HumanOperators
Adam Marcus (MIT CSAIL), Eugene Wu (MIT CSAIL), David R. Karger (MIT CSAIL), Samuel Madden (MIT CSAIL), Robert C. Miller (MIT CSAIL)

Automatic Example Queries for Ad Hoc Databases
Bill Howe (University of Washington), Garret Cole (University of Washington), Nodira Khoussainova (University of Washington), Leilani Battle (University of Washington)

NetTrails: A Declarative Platform for Maintaining and Querying Provenance in Distributed Systems
Wenchao Zhou (University of Pennsylvania), Qiong Fei (University of Pennsylvania), Shengzhi Sun (University of Pennsylvania), Tao Tao (University of Pennsylvania), Andreas Haeberlen (University of Pennsylvania), Zachary Ives (University of Pennsylvania), Boon Thau Loo (University of Pennsylvania), Micah Sherr (Georgetown University)

GBLENDER: Visual Subgraph Query Formulation Meets Query Processing
Changjiu Jin (Nanyang Technological University), Sourav S Bhowmick (Nanyang Technological University), Xiaokui Xiao (Nanyang Technological University), Byron Choi (Hong Kong Baptist University), Shuigeng Zhou (Fudan University)

Coordination through Querying in the Youtopia System
Nitin Gupta (Cornell University), Lucja Kot (Cornell University), Gabriel Bender (Cornell University), Sudip Roy (Cornell University), Johannes Gehrke (Cornell University), Christoph Koch (EPFL)

DBWiki: A Structured Wiki for Curated Data and Collaborative Data Management
Peter Buneman (University of Edinburgh), James Cheney (University of Edinburgh), Sam Lindley (University of Edinburgh), Heiko Müller (CSIRO)

Rapid Development of Web-Based Query Interfacesfor XML Datasets with QURSED
Abhijith Kashyap (SUNY at Buffalo), Michalis Petropoulos (SUNY at Buffalo)

WEDNESDAY 15TH 10:30 – 12:00

SIGMOD Research 13: Streams and Complex Event Processing
Location: Olympia A

Changing Flights in Mid-air: A Model for Safely Modifying Continuous Queries
Kyumars Sheykh Esmaili (Systems Group), Tahmineh Sanamrad (ETH Zurich), Peter M Fischer (ETH Zurich), Nesime Tatbul (ETH Zurich)

How Soccer Players Would Do Stream Joins
Jens Teubner (ETH Zurich), Rene Mueller (IBM Almaden)

BE-Tree: An Index Structure to Efficiently Match Boolean Expressions over High-dimensional Discrete Space
Mohammad Sadoghi (University of Toronto), Hans-Arno Jacobsen (University of Toronto)

TI: An Efficient Indexing Mechanism for Real-Time Search on Tweets
Chun Chen (Zhejiang University), Feng Li (National University of Singapore), Beng Chin Ooi (National University of Singapore), Sai Wu (National University of Singapore)

SIGMOD Research 14: Query Processing
Location: Olympia B

More Efficient Datalog Queries: Subsumptive Tabling Beats Magic Sets
K. Tuncay Tekle (LogicBlox), Yanhong A. Liu (State University of New York at Stony Brook)

Entangled Queries: Enabling Declarative Data-Driven Coordination
Nitin Gupta (Cornell University), Lucja Kot (Cornell University), Sudip Roy (Cornell University), Gabriel Bender (Cornell University), Johannes Gehrke (Cornell University), Christoph Koch (EPFL)

Data Generation using Declarative Constraints
Arvind Arasu (Microsoft Research), Raghav Kaushik (Microsoft Research), Jian Li (University of Maryland)

Efficient Auditing For Complex SQL queries
Raghav Kaushik (Microsoft Research), Ravi Ramamurthy (Microsoft Corporation)

SIGMOD Research 15: Data Mining
Location: Vergina

Exact Indexing for Support Vector Machines
Hwanjo Yu (POSTECH (Pohang University of Science and Technology)), Ilhwan Ko (POSTECH (Pohang University of Science and Technology)), Youngdae Kim (POSTECH (Pohang University of Science and Technology)), Seungwon Hwang (POSTECH (Pohang University of Science and Technology)), Wook-Shin Han (Kyungpook National University)

Local Graph Sparsification for Scalable Clustering
Venu Satuluri (The Ohio State University), Srinivasan Parthasarathy (The Ohio State University), Yiye Ruan (The Ohio State University)

Advancing Data Clustering via Projective Clustering Ensembles
Francesco Gullo (University of Calabria), Carlotta Domeniconi (George Mason University), Andrea Tagarelli (University of Calabria)

Sampling Based Algorithms for Quantile Computation in Sensor Networks
Zengfeng Huang (Hong Kong University of Science and Technology), Lu Wang (Hong Kong University of Science and Technology), Ke Yi (Hong Kong University of Science and Technology), Yunhao Liu (Hong Kong University of Science and Technology)

SIGMOD Industrial Keynote and Session 4: Business Analytics
Location: Pella

SIGMOD Industrial Keynote 2: Technology and Trends for Smarter Business Analytics
Don Campbell (IBM)

LCI: A Social Channel Analysis Platform for Live Customer Intelligence
Malu Castellanos (HP), Umeshwar Dayal (HP), Meichun Hsu (HP), Riddhiman Ghosh (HP), Mohamed Dekhil (HP), Yue Lu (UIUC), Lei Zhang (University of Illinois at Chicago)

SIGMOD Tutorial 4: Web Data Management
Location: Athens View
Presenters: Michael J Cafarella (University of Michigan), Alon Y Halevy (Google, Inc.)

PODS Research 6: Semistructured Data and XML
Location: Horizon

The Complexity of Text-Preserving XML Transformations
Timos Antonopoulos (Hasselt University and Transnational University of Limburg), Wim Martens (TU Dortmund), Frank Neven (Hasselt University and Transnational University of Limburg)

Efficient Evaluation for a Temporal Logic on Changing XML Documents
Mikolaj Bojańczyk (University of Warsaw), Diego Figueira (University of Warsaw & University of Edinburgh)

Finding a Minimal Tree Pattern Under Neighborhood Constraints

Benny Kimelfeld (IBM Research - Almaden), Yehoshua Sagiv (The Hebrew University)

SIGMOD Demo B & Programming Contest: System and Performance
Location: Mycenae

One-pass Data Mining Algorithms in a DBMS with UDFs
Carlos Ordonez (University of Houston), Sasi K Pitchaimalai (University of Houston)

Inspector Gadget: A Framework for Custom Monitoring and Debugging of Distributed Dataflows
Christopher Olston (Yahoo! Research), Benjamin Reed (Yahoo! Research)

RAFT at Work: Speeding-Up MapReduce Applications under Task and Node Failures
Jorge-Arnulfo Quiané-Ruiz (Saarland University), Christoph Pinkel (Saarland University), Jörg Schad (Saarland University), Jens Dittrich (Saarland University)

WattDB: An Energy-Proportional Cluster of Wimpy Nodes
Daniel Schall (TU Kaiserslautern), Volker Hudlet (TU Kaiserslautern)

BRRL: A Recovery Library for Main-Memory Applications in the Cloud
Tuan Cao (Cornell University), Benjamin Sowell (Cornell University), Marcos Vaz Salles (Cornell University), Alan Demers (Cornell University), Johannes Gehrke (Cornell University)

A Data-oriented Transaction Execution Engine and Supporting Tools
Ippokratis Pandis (Carnegie Mellon University), Pinar Tozun (Ecole Polytechnique Federale de Lausanne), Miguel Branco (Ecole Polytechnique Federale de Lausanne), Dimitris Karampinas (University of Patras), Danica Porobic (Ecole Polytechnique Federale de Lausanne), Ryan Johnson (University of Toronto), Anastasia Ailamaki (Ecole Polytechnique Federale de Lausanne)

iGraph in Action: Performance Analysis of Disk-Based Graph Indexing Techniques
Wook-Shin Han (Kyungpook National University), Minh-Duc Pham (Kyungpook National University), Jinsoo Lee (Kyungpook National University), Romans Kasperovics (Kyungpook National University), Jeffrey Xu Yu (Chinese University of Hong Kong)

StreamRec: A Real-Time Recommender System
Badrish Chandramouli (Microsoft Research), Justin J Levandoski (University of Minnesota), Ahmed Eldawy (University of Minnesota), Mohamed F Mokbel (University of Minnesota)

WEDNESDAY 15TH 16:00 – 18:00

SIGMOD Plenary I: Posters of Research Papers from Sessions 11-21
Location: Olympia

PODS Research 7: Rule-based Query Languages
Location: Horizon

A Rule-based Language for Web Data Management
Serge Abiteboul (INRIA Saclay & LSV-ENS Cachan), Meghyn Bienvenu (CNRS & Université Paris Sud), Alban Galland, Émilien Antoine (INRIA Saclay & LSV-ENS Cachan)

Relational Transducers for Declarative Networking
Tom J. Ameloot, Frank Neven, Jan Van den Bussche (Hasselt University & Transnational University of Limburg)

Rewrite Rules for Search Database Systems
Ronald Fagin, Benny Kimelfeld, Yunyao Li (IBM Research - Almaden), Sriram Raghavan (IBM India Research Lab), Shivakumar Vaithyanathan (IBM Research - Almaden)

SIGMOD Demo A: Ranking, the Web and Social Media
Location: Mycenae

SkylineSearch: Semantic Ranking and Result Visualization for PubMed
Julia Stoyanovich (University of Pennsylvania), Mayur Lodha (Columbia University), William Mee (Columbia University), Kenneth A. Ross (Columbia University)

A Cross-Service Travel Engine for Trip Planning
Gang Chen (Zhejiang University), Chen Liu (National University of Singapore), Meiyu Lu (National University of Singapore), Beng Chin Ooi (National University of Singapore), Shanshan Ying (National University of Singapore), Anthony Tung (National University of Singapore), Dongxiang Zhang (National University of Singapore), Meihui Zhang (National University of Singapore)

WINACS: Construction and Analysis of Web-Based Computer Science Information Networks
Tim Weninger (University of Illinois Urbana-Champaign), Marina Danilevsky (University of Illinois Urbana-Champaign), Fabio Fumarola (Universita), Joshua Hailpern (University of Illinois Urbana-Champaign), Jiawei Han (University of Illinois Urbana-Champaign), Thomas J. Johnston (University of Illinois Urbana-Champaign), Surya Kallumadi (Kansas State University), Hyungsul Kim (University of Illinois Urbana-Champaign), Zhijin Li (University of Illinois Urbana-Champaign), David McCloskey (University of Illinois Urbana-Champaign), Yizhou Sun (University of Illinois Urbana-Champaign), Nathan E. TeGrotenhuis (Whitworth University), Chi Wang (University of Illinois Urbana-Champaign), Xiao Yu (University of Illinois Urbana-Champaign)

Tweets as Data: Demonstration of TweeQL and TwitInfo
Adam Marcus (MIT CSAIL), Michael S. Bernstein (MIT CSAIL), Osama Badar (MIT CSAIL), David R. Karger (MIT CSAIL), Samuel Madden (MIT CSAIL), Robert C. Miller (rcm@csail.mit.edu)

MOBIES: Mobile-Interface Enhancement Service for Hidden Web Database
Xin Jin (George Washington University), Aditya Mone (University of Texas at Arlington), Nan Zhang (George Washington University), Gautam Das (University of Texas at Arlington)

Search Computing: Multi-domain Search on Ranked Data
Alessandro Bozzon (Politecnico di Milano), Daniele Braga (Politecnico di Milano), Marco Brambilla (Politecnico di Milano), Stefano Ceri (Politecnico di Milano), Francesco Corcoglioniti (Politecnico di Milano), Piero Fraternali (Politecnico di Milano), Salvatore Vadacca (Politecnico di Milano)

EnBlogue -- Emergent Topic Detection in Web 2.0 Streams
Foteini Alvanaki (Saarland University), Michel Sebastian (Saarland University), Krithi Ramamritham (IIT Bombay), Gerhard Weikum (Max-Planck Institute Informatics)

NOAM: News Outlets Analysis and Monitoring System
Ilias Flaounas (University of Bristol), Omar Ali (University of Bristol), Marco Turchi (European Commission), Tristan Snowsill (University of Bristol), Florent Nicart (Université de Rouen), Tijl De Bie (University of Bristol), Nello Cristianini (University of Bristol)

THURSDAY 16TH 09:00 – 10:00

SIGMOD Keynote 2: Managing Scientific Data
Location: Olympia
Speaker: Anastasia Ailamaki (EPFL)

THURSDAY 16TH 10:30 – 12:00

SIGMOD Research 16: Information Retrieval
Location: Olympia A

Context-sensitive Ranking for Document Retrieval
Liang Jeff Chen (UC San Diego), Yannis Papakonstantinou (UC San Diego)

Score-Consistent Algebraic Optimization of Full-Text Search Queries with GRAFT
Nathan Bales (UC San Diego), Alin Deutsch (UC San Diego), Vasilis Vassalos (A.U.E.B)

Efficient Diversity-Aware Search
Albert Angel (University of Toronto), Nick Koudas (University of Toronto)

Mining a Search Engine's Corpus: Efficient Yet Unbiased Sampling and Aggregate Estimation
Mingyang Zhang (George Washington University), Nan Zhang (George Washington University), Gautam Das (University of Texas at Arlington)

SIGMOD Research 17: Probabilistic and Uncertain Databases
Location: Olympia B

Ranking with Uncertain Scoring Functions: Semantics and Sensitivity Measures
Mohamed A Soliman (Greenplum), Ihab F Ilyas (University of Waterloo), Davide Martinenghi (Politecnico di Milano), Marco Tagliasacchi (Politecnico di Milano)

Querying Uncertain Data with Aggregate Constraints

Mohan Yang (Shanghai Jiao Tong University), Haixun Wang (Microsoft Research Asia), Haiquan Chen (Auburn University), Wei-Shinn Ku (Auburn University)

Jigsaw: Efficient optimization over uncertain enterprise data
Oliver A Kennedy (EPFL), Suman Nath (Microsoft Research)

Sensitivity Analysis and Explanations for Robust Query Evaluation in Probabilistic Databases
Bhargav Kanagal (University of Maryland), Jian Li (University of Maryland), Amol Deshpande (University of Maryland)

SIGMOD Research 18: Scalable Data Analytics
Location: Vergina

Fast Personalized PageRank on MapReduce
Bahman Bahmani (Stanford University), Kaushik Chakrabarti (Microsoft Research), Dong Xin (Google Inc.)

Processing Theta-Joins using MapReduce
Alper Okcan (Northeastern University), Mirek Riedewald (Northeastern University)

Llama: Leveraging Columnar Storage for Scalable Join Processing in the MapReduce Framework
Yuting Lin (National University of Singapore), Divyakant Agrawal (University of California, Santa Barbara), Chun Chen (Zhejiang University), Beng Chin Ooi (National University of Singapore), Sai Wu (National University of Singapore)

A Platform for Scalable One-Pass Analytics using MapReduce
Boduo Li (University of Massachusetts Amherst), Edward Mazur (University of Massachusetts Amherst), Yanlei Diao (University of Massachusetts Amherst), Andrew McGregor (University of Massachusetts Amherst), Prashant Shenoy (University of Massachusetts Amherst)

SIGMOD Tutorial 5: Large-Scale Copy Detection
Location: Athens View
Presenters: Xin Luna Dong (AT&T Labs - Research), Divesh Srivastava (AT&T Labs - Research)

SIGMOD Panel 1: Data Management Issues in Health and Medical Informatics
Location: Horizon
Moderator: Tamer Özsu (University of Waterloo)

SIGMOD Demo D: User Support and Development Environments
Location: Mycenae

CONFLuEnCE: CONtinuous workFLow ExeCution Engine
Panayiotis Neophytou (University of Pittsburgh), Panos K. Chrysanthis (University of Pittsburgh), Alexandros Labrinidis (University of Pittsburgh)

Demonstration of Qurk: A Query Processor for HumanOperators
Adam Marcus (MIT CSAIL), Eugene Wu (MIT CSAIL), David R. Karger (MIT CSAIL), Samuel Madden (MIT CSAIL), Robert C. Miller (MIT CSAIL)

Automatic Example Queries for Ad Hoc Databases
Bill Howe (University of Washington), Garret Cole (University of Washington), Nodira Khoussainova (University of Washington), Leilani Battle (University of Washington)

NetTrails: A Declarative Platform for Maintaining and Querying Provenance in Distributed Systems
Wenchao Zhou (University of Pennsylvania), Qiong Fei (University of Pennsylvania), Shengzhi Sun (University of Pennsylvania), Tao Tao (University of Pennsylvania), Andreas Haeberlen (University of Pennsylvania), Zachary Ives (University of Pennsylvania), Boon Thau Loo (University of Pennsylvania), Micah Sherr (Georgetown University)

GBLENDER: Visual Subgraph Query Formulation Meets Query Processing
Changjiu Jin (Nanyang Technological University), Sourav S Bhowmick (Nanyang Technological University), Xiaokui Xiao (Nanyang Technological University), Byron Choi (Hong Kong Baptist University), Shuigeng Zhou (Fudan University)

Coordination through Querying in the Youtopia System
Nitin Gupta (Cornell University), Lucja Kot (Cornell University), Gabriel Bender (Cornell University), Sudip Roy (Cornell University), Johannes Gehrke (Cornell University), Christoph Koch (EPFL)

DBWiki: A Structured Wiki for Curated Data and Collaborative Data Management

Peter Buneman (University of Edinburgh), James Cheney (University of Edinburgh), Sam Lindley (University of Edinburgh), Heiko Müller (CSIRO)

Rapid Development of Web-Based Query Interfacesfor XML Datasets with QURSED
Abhijith Kashyap (SUNY at Buffalo), Michalis Petropoulos (SUNY at Buffalo)

THURSDAY 16TH 13:30 – 15:00

SIGMOD Research 19: Graph Management
Location: Olympia A

Neighborhood Based Fast Graph Search in Large Networks
Arijit Khan (University of California, Santa Barbara), Nan Li (University of California, Santa Barbara), Xifeng Yan (University of California, Santa Barbara), Ziyu Guan (University of California, Santa Barbara), Supriyo Chakraborty (University of California, Los Angeles), Shu Tao (IBM T. J. Watson Research Center)

A Memory Efficient Reachability Data Structure Through Bit Vector Compression
Sebastiaan J. van Schaik (University of Oxford & Utrecht University), Oege de Moor (University of Oxford)

Incremental Graph Pattern Matching
Wenfei Fan (University of Edinburgh), Jianzhong Li (Harbin Institute of Technologies), Jizhou Luo (Harbin Institute of Technologies), Zijing Tan (Fudan University), Xin Wang (University of Edinburgh), Yinghui Wu (University of Edinburgh)

Assessing and Ranking Structural Correlations in Graphs
Ziyu Guan (University of California), Jian Wu (Zhejiang University), Qing Zhang (TaoBao.com), Ambuj Singh (University of California), Xifeng Yan (University of California)

SIGMOD Research 20: OLAP and its Applications
Location: Olympia B

Graph Cube: On Warehousing and OLAP Multidimensional Networks
Peixiang Zhao (University of Illinois at Urbana-Champaign), Xiaolei Li (Microsoft Corporation), Dong Xin (Google Inc.), Jiawei Han (University of Illinois at Urbana-Champaign)

MaSM: Efficient Online Updates in Data Warehouses
Manos Athanassoulis (Ecole Polytechnique Federale de Lausanne), Shimin Chen (Intel Labs), Anastasia Ailamaki (Ecole Polytechnique Federale de Lausanne), Phillip B. Gibbons (Intel Labs), Radu Stoica (Ecole Polytechnique Federale de Lausanne)

Latent OLAP: Data Cubes over Latent Variables
Deepak Agarwal (Yahoo! Research), Bee-Chung Chen (Yahoo! Research)

E-Cube: Multi-Dimensional Event Sequence Analysis Using Hierarchical Pattern Query Sharing
Mo Liu (Worcester Polytechnic Institute), Elke Rundensteiner (Worcester Polytechnic Institute), Kara Greenfield (Worcester Polytechnic Institute), Chetan Gupta (Hewlett-Packard Labs), Song Wang (Hewlett-Packard Labs), Ismail Ari (Ozyegin University), Abhay Mehta (Hewlett-Packard Labs)

SIGMOD Research 21: Similarity Search and Queries
Location: Pella

ATLAS: A Probabilistic Algorithm for High Dimensional Similarity Search
Jiaqi Zhai (Cornell University), Yin Lou (Cornell University), Johannes Gehrke (Cornell University)

Flexible Aggregate Similarity Search
Yang Li (Shanghai Jiao Tong University), Feifei Li (Florida State University), Ke Yi (Hong Kong University of Science and Technology), Bin Yao (Florida State University), Min Wang (HP Labs China)

Effective Data Co-Reduction for Multimedia Similarity Search
Zi Huang (The University of Queensland), Hengtao Shen (The University of Queensland), Jiajun Liu (The University of Queensland), Xiaofang Zhou (The University of Queensland)

Efficient Exact Edit Similarity Query Processing with the Asymmetric Signature Scheme
Jianbin Qin (University of New South Wales), Wei Wang (University of New South Wales), Yifei Lu (University of New South Wales), Chuan Xiao (University of New South Wales), Xuemin Lin (University of New South Wales & East Normal China University)

SIGMOD Tutorial 6: Data Management Over Flash Memory
Location: Athens View

Presenters: Ioannis Koltsidas (IBM Research), Stratis D Viglas (University of Edinburgh)

SIGMOD Industrial 5: Dynamic Optimization and Unstructured Content
Location: Pella

Turbocharging DBMS Buffer Pool Using SSDs
Jaeyoung Do (University of Wisconsin-Madison), Donghui Zhang (Microsoft Jim Gray Systems Lab), Jignesh M Patel (University of Wisconsin-Madison), David J DeWitt (Microsoft Jim Gray Systems Lab), Jeffrey F Naughton (University of Wisconsin-Madison), Alan Halverson (Microsoft Jim Gray Systems Lab)

Online Reorganization in Read Optimized MMDBS
Felix Beier (Ilmenau University of Technology), Knut Stolze (IBM Research Development), Kai-Uwe Sattler (Ilmenau University of Technology)

Automated Partitioning Design in Parallel Database Systems
Rimma Nehme (Microsoft Jim Gray Systems Lab), Nicolas Bruno (Microsoft)

Oracle Database Filesystem
Krishna Kunchithapadam (Oracle Corporation), Wei Zhang (Oracle Corporation), Amit Ganesh (Oracle Corporation), Niloy Mukherjee (Oracle Corporation)

SIGMOD Demo C & Programming Contest Winners: Data Integration and Probabilistic Databases
Location: Mycenae

Pay-As-You-Go Mapping Selection in Dataspaces
Cornelia Hedeler (The University of Manchester), Khalid Belhajjame (The University of Manchester), Norman W Paton (The University of Manchester), Alvaro A A Fernandes (The University of Manchester), Suzanne M Embury (The University of Manchester), Lu Mao (The University of Manchester), Chenjuan Guo (The University of Manchester)

Exelixis: Evolving Ontology-Based Data Integration System
Haridimos Kondylakis (FORTH-ICS), Dimitris Plexousakis (FORTH-ICS)

U-MAP: A System for Usage-Based Schema Matching and Mapping
Hazem Elmeleegy (AT&T Labs - Research), Jaewoo Lee (Purdue University), El Kindi Rezig (Purdue University), Mourad Ouzzani (Purdue University), Ahmed Elmagarmid (Qatar Computing Research Institute, Qatar Foundation)

The SystemT IDE: An Integrated Development Environment for Information Extraction Rules
Laura Chiticariu (IBM Research - Almaden), Vivian Chu (IBM research - Almaden), Sajib Dasgupta (IBM Research - Almaden), Thilo W Goetz (IBM Software - Germany), Howard Ho (IBM Research - Almaden), Rajasekar Krishnamurthy (IBM Research - Almaden), Alexander Lang (IBM Software - Germany), Yunyao Li (IBM Research - Almaden), Bin Liu (University of Michigan), Sriram Raghavan (IBM Research - India), Frederick R Reiss (IBM Research - Almaden), Shivakumar Vaithyanathan (IBM Research - Almaden), Huaiyu Zhu (IBM Research - Almaden)

ProApproX: A Lightweight Approximation Query Processor over Probabilistic Trees
Pierre Senellart (Institut Télécom; Télécom ParisTech), Asma Souihli (Institut Télécom; Télécom ParisTech)

SPROUT²: A Squared Query Engine for Uncertain Web Data
Robert Fink (University of Oxford), Andrew Hogue (Google Inc.), Dan Olteanu (University of Oxford), Swaroop Rath (University of Oxford)

Fuzzy Prophet: Parameter exploration in uncertain enterprise scenarios
Oliver A Kennedy (EPFL), Steve Lee (Microsoft Corporation), Charles Loboz (Microsoft Corporation), Slawek Smyl (Microsoft Corporation), Suman Nath (Microsoft Research)

LinkDB: A Probabilistic Linkage Database System
Ekaterini Ioannou (Technical University of Crete), Wolfgang Nejdl (L3S Research Center), Claudia Niederée (L3S Research Center), Yannis Velegrakis (University of Trento)

THURSDAY 16TH 15:30 – 17:30

SIGMOD Plenary II: Posters of Research Papers from Sessions 1-10
Location: Olympia